

Data Science

Final Project No 7

House Sales in King Country, USA

In this project, I have to perform as Data Analyst working in a Real Estate Investment Trust. They'd like starting investing in Residential real estate. I'll determine the market price of a house given a set of features and predict housing prices

Élaboré par :
Smide ALEXIS

Encadré par :
Coursera Staff

Data Science with Python

15 Janvier 2024

Importation des packages

```
import pandas as pd
import seaborn as sbn
import numpy as np
import scipy as sc
import scipy.stats as stats
import matplotlib.pyplot as plt
import pipeline
import tqdm
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
import warnings
warnings.filterwarnings("ignore", category=UserWarning)
```

Importation du dataset

```
filepath=(
    'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkill'
    'sNetwork-DA0101EN-SkillsNetwork/labs/FinalModule_Coursera/data/kc_house_data_NaN.csv'
)
df = pd.read_csv(filepath, header=0)
df.head()
```

	Unnamed: 0	id	date	...	long	sqft_living15	sqft_lot15
0	0	7129300520	20141013T000000	...	-122.257	1340	5650
1	1	6414100192	20141209T000000	...	-122.319	1690	7639
2	2	5631500400	20150225T000000	...	-122.233	2720	8062
3	3	2487200875	20141209T000000	...	-122.393	1360	5000
4	4	1954400510	20150218T000000	...	-122.045	1800	7503

[5 rows x 22 columns]

```
df.tail()
```

	Unnamed: 0	id	...	sqft_living15	sqft_lot15
21608	21608	263000018	...	1530	1509
21609	21609	6600060120	...	1830	7200
21610	21610	1523300141	...	1020	2007
21611	21611	291310100	...	1410	1287
21612	21612	1523300157	...	1020	1357

[5 rows x 22 columns]

Display the data types of each column using the function dtypes.

```
df.dtypes
```

```
Unnamed: 0      int64
id              int64
date            object
price           float64
bedrooms        float64
bathrooms       float64
sqft_living      int64
sqft_lot         int64
floors          float64
waterfront      int64
view            int64
condition       int64
grade           int64
sqft_above      int64
sqft_basement   int64
yr_built        int64
yr_renovated    int64
zipcode         int64
lat             float64
long            float64
sqft_living15   int64
sqft_lot15      int64
dtype: object
```

We use the method describe to obtain a statistical summary of the dataframe.

```
df.describe()
```

	Unnamed: 0	id	...	sqft_living15	sqft_lot15
count	21613.000000	2.161300e+04	...	21613.000000	21613.000000
mean	10806.000000	4.580302e+09	...	1986.552492	12768.455652
std	6239.28002	2.876566e+09	...	685.391304	27304.179631
min	0.00000	1.000102e+06	...	399.000000	651.000000
25%	5403.000000	2.123049e+09	...	1490.000000	5100.000000
50%	10806.000000	3.904930e+09	...	1840.000000	7620.000000
75%	16209.000000	7.308900e+09	...	2360.000000	10083.000000
max	21612.000000	9.900000e+09	...	6210.000000	871200.000000

```
[8 rows x 21 columns]
```

Drop the columns “id” and “Unnamed : 0” from axis 1 using the method drop(), then use the method describe() to obtain a statistical summary of the data. Make sure the inplace parameter is set to True.

```
df.drop(["id","Unnamed: 0"], axis=1, inplace=True)
df.describe()
```

	price	bedrooms	...	sqft_living15	sqft_lot15
count	2.161300e+04	21600.000000	...	21613.000000	21613.000000
mean	5.400881e+05	3.372870	...	1986.552492	12768.455652

```

std      3.671272e+05      0.926657 ...      685.391304      27304.179631
min      7.500000e+04      1.000000 ...      399.000000      651.000000
25%      3.219500e+05      3.000000 ...      1490.000000      5100.000000
50%      4.500000e+05      3.000000 ...      1840.000000      7620.000000
75%      6.450000e+05      4.000000 ...      2360.000000      10083.000000
max      7.700000e+06      33.000000 ...      6210.000000      871200.000000

```

```
[8 rows x 19 columns]
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  21613 non-null  object
1   price                 21613 non-null  float64
2   bedrooms              21600 non-null  float64
3   bathrooms              21603 non-null  float64
4   sqft_living           21613 non-null  int64
5   sqft_lot              21613 non-null  int64
6   floors                21613 non-null  float64
7   waterfront            21613 non-null  int64
8   view                  21613 non-null  int64
9   condition             21613 non-null  int64
10  grade                 21613 non-null  int64
11  sqft_above            21613 non-null  int64
12  sqft_basement         21613 non-null  int64
13  yr_built              21613 non-null  int64
14  yr_renovated          21613 non-null  int64
15  zipcode               21613 non-null  int64
16  lat                   21613 non-null  float64
17  long                  21613 non-null  float64
18  sqft_living15         21613 non-null  int64
19  sqft_lot15            21613 non-null  int64
dtypes: float64(6), int64(13), object(1)
memory usage: 3.3+ MB

```

We can see we have missing values for the columns bedrooms and bathrooms

```
print("number of NaN values for the column bedrooms:", df['bedrooms'].isnull().sum())
```

```
number of NaN values for the column bedrooms: 13
```

```
print("number of NaN values for the column bathrooms:", df['bathrooms'].isnull().sum())
```

```
number of NaN values for the column bathrooms: 10
```

```
mean=df['bedrooms'].mean()
df['bedrooms'].replace(np.nan,mean, inplace=True)

mean=df['bathrooms'].mean()
df['bathrooms'].replace(np.nan,mean, inplace=True)
```


Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a data frame.

```
unique_floor_values = df["floors"].value_counts().to_frame()
unique_floor_values
```

	count
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

Use the function `boxplot` in the `seaborn` library to determine whether houses with a waterfront view or without a waterfront view have more price outliers.


```
sbn.boxplot(x="waterfront", y="price", data=df)
```



Use the function `regplot` in the `seaborn` library to determine if the feature `sqft_above` is negatively or positively correlated with `price`.

```
sbn.regplot(x="sqft_above", y="price", data=df, line_kws={"color": "red"})  
plt.ylim(0,)
```

```
(0.0, 8081250.0)
```



We can use the Pandas method `corr()` to find the feature other than `price` that is most correlated with `price`.

```
# df.corr()['price'].sort_values()
```

We can Fit a linear regression model using the longitude feature `'long'` and calculate the R^2 .

```
X = df[["long"]]  
Y = df[["price"]]  
lm = LinearRegression()  
lm.fit(X,Y)
```

```
LinearRegression()
```

```
print(lm.score(X,Y))
```

```
0.00046769430149007363
```

Fit a linear regression model to predict the 'price' using the feature 'sqft_living' then calculate the R^2 .

```
V = df[["sqft_living"]]
Y = df[["price"]]
lm = LinearRegression()
lm.fit(V,Y)
```

```
LinearRegression()
```

```
print(lm.score(V,Y))
```

```
0.4928532179037931
```

Fit a linear regression model to predict the 'price' using the list of features.

```
features =[
    "floors","waterfront","lat","bedrooms","sqft_basement","view",
    "bathrooms","sqft_living15","sqft_above","grade","sqft_living"
]
Z = df[features]
lm.fit(Z,Y)
```

```
LinearRegression()
```

```
print(lm.score(Z, Y))
```

```
0.6576868690521125
```

Pipeline

Create a list of tuples, the first element in the tuple contains the name of the estimator : 'scale', 'polynomial', 'model'. The second element in the tuple contains the model constructor StandardScaler(), PolynomialFeatures(include_bias=False), LinearRegression()

```
Input=[
    ('scale',StandardScaler()), ('polynomial', PolynomialFeatures(include_bias=False)),
    ('model', LinearRegression())
]
```

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list features, and calculate the R^2 .

```
pipeline=Pipeline(Input)
Z = Z.astype(float)
pipeline.fit(Z,Y)
```

```
Pipeline(steps=[('scale', StandardScaler()),
                 ('polynomial', PolynomialFeatures(include_bias=False)),
                 ('model', LinearRegression())])
```

```
ypipe=pipeline.predict(Z)
print(r2_score(Y,ypipe))
```

0.751335953931385

We will split the data into training and testing sets :

```
features =[
    "floors","waterfront","lat","bedrooms","sqft_basement","view",
    "bathrooms","sqft_living15","sqft_above","grade","sqft_living"
]
X = df[features]
Y = df['price']

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.15, random_state=1)

print("number of test samples:", x_test.shape[0])
```

number of test samples: 3242

```
print("number of training samples:",x_train.shape[0])
```

number of training samples: 18371

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data.

```
RidgeModel=Ridge(alpha=0.1)
RidgeModel.fit(x_train, y_train)
```

```
Ridge(alpha=0.1)
```

```
yhat = RidgeModel.predict(x_test)
print(r2_score(y_test,yhat))
```

0.6478759163939111

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided.

```
pr = PolynomialFeatures(degree=2)
x_train_pr = pr.fit_transform(x_train)
x_test_pr = pr.fit_transform(x_test)
RidgeModel.fit(x_train_pr, y_train)
```

```
Ridge(alpha=0.1)
```

```
y_hat = RidgeModel.predict(x_test_pr)
print(r2_score(y_test,y_hat))
```

```
0.7002744267811738
```